

# Agustin Leon

al8937@nyu.edu | 929-261-9361 | [website/portfolio](#) | [linkedin](#) | [github](#) | Brooklyn, NY

## EDUCATION

### New York University, Tandon School of Engineering

New York, USA

*Master of Science in Computer Engineering*

*Sep 2024 – May 2026*

- GPA: 4.0/4.0
- **Graduate Teaching Assistant** for foundations of computer science course including algorithm design, data structures, operating systems, and C++ programming, mentoring cohorts of over 150 students.
- **Courses:** Machine Learning, MLOps, Deep Learning, Foundations of Robotics, Embedded Systems, Computing Systems Architecture. **Upcoming:** Computer Vision, Reinforcement Learning, Transformers-NLP-LLMs.

### University of Chile

Santiago, Chile

*Bachelor of Science in Economics (distinction)*

*Mar 2013 – Dec 2017*

## PROJECTS

### ‘Spotify Buddies’ – Full ML Pipeline for Playlist Recommender | *MLOps, Docker, Ray*

Spring 2025

- Full-stack ML system to recommend organic user-created playlists based on musical taste similarity, using dual-model architecture (user pairing via KNN + playlist ranking via Bayesian Personalized Ranking).
- Led training and experiment tracking with MLflow, and hyperparameter optimization using Ray.
- Integrated the system with a CI/CD pipeline on Kubernetes using Docker, Prometheus, and FastAPI for serving, load testing, and degradation monitoring.

### LoRA Fine-Tuning with Domain-Adaptive Data Augmentation | *HuggingFace, PEFT*

Spring 2025

- Fine-tuned a RoBERTa-base model using Low-Rank Adaptation (LoRA), reducing trainable parameters to 917K.
- Improved model robustness with domain-adaptive data augmentation by fine-tuning GPT-2 to generate 12k synthetic text samples from unlabeled data using custom randomized prompts to prevent overfitting.
- Ranked 1st out of 400 participants in a [Kaggle competition](#), with 93.8% validation and 86.95% test accuracy on a held-out OOD test dataset

### Image classification with parameter-constrained ResNets | *PyTorch, Data Augmentation*

Fall 2024

- Trained a light ResNet with less than 5-million parameters on the CIFAR-10 dataset.
- Achieved 96% test accuracy on held-out test data, applying diverse techniques for data augmentation.

## EXPERIENCE

### Machine Learning Engineer Intern

New York, USA

NEVARA

*May 2025 – Aug 2025*

- Selected to lead a team of 5 interns, setting priorities, coordinating tasks, removing blockers, and ensuring we delivered 4 out of 4 product milestones successfully and on time.
- Architected and deployed an AI-powered sales coaching platform on AWS, integrating FastAPI microservices with Supabase. Developed LLM pipelines with prompt engineering, input compression, retrieval-augmented generation (RAG), lightweight agent patterns, and a feedback loop for continuous improvement from user feedback.
- Reduced sales coach workload by 90% (from 10 hours of manual call review and report writing, to 1 hour) through automated note generation and multi-call report synthesis.

### Associate Research Consultant

Santiago, Chile

GBA ECONOMIC CONSULTING

*Jan 2022 – Sep 2024*

- Conducted quantitative research for 15 projects across industries (energy, telecom, transportation, IP, antitrust).
- Delivered frequent presentations to diverse audiences, including C-suite executives and non-technical teams.
- Drove \$10M in client savings and profit gains; key clients included Mastercard, Engie, and LATAM Airlines.

### Data Scientist

Santiago, Chile

MINISTRY OF EDUCATION, GOVERNMENT OF CHILE

*Jul 2019 – Dec 2021*

- Led a team to develop a student dropout classification model, increasing accuracy from 50% to 80%.
- Model gave insights for policy aimed at mitigating increased school dropout rates during the pandemic.

**Technical Skills:** Languages: Python, C/C++, JavaScript, R | **ML/DL:** PyTorch, TensorFlow, HuggingFace, scikit-learn, Keras, ONNX, LLMs (prompt engineering, RAG, agents) | **Web:** REST APIs, Node.js, FastAPI | **Tools:** Git, Docker, AWS, GCP, Kubernetes, MLflow, CI/CD | **Data:** Numpy, Matplotlib, MongoDB, Supabase, MATLAB | **Robotics:** ROS2